

Indian Institute of Technology Gandhinagar



Assignment 3:

Pretraining and fine-tuning an LLM

CS 613: NLP

Group T10: Exquisiters

Aditi Dey (20110007)

Anavart Pandya (20110016)

Ayush Gupta (20110031)

Mann Kumar Jain (20110108)

Pinki Kumari(22210028)

Ronak Kalra (20110171)

Sakshi Jain (20110181)

Vedang Chavan (20110222)

Task 1: Bert-base-uncased model.

Task 2: Parameters of the selected model from the code

| Modules | Parameters |
|---|------------|
| embeddings.word_embeddings.weight | 23440896 |
| embeddings.position_embeddings.weight | 393216 |
| embeddings.token_type_embeddings.weight | 1536 |
| embeddings.LayerNorm.weight | 768 |
| embeddings.LayerNorm.bias | 768 |
| encoder.layer.0.attention.self.query.weight | 589824 |
| encoder.layer.0.attention.self.query.bias | 768 |
| encoder.layer.0.attention.self.key.weight | 589824 |
| encoder.layer.0.attention.self.key.bias | 768 |
| encoder.layer.0.attention.self.value.weight | 589824 |
| encoder.layer.0.attention.self.value.bias | 768 |
| encoder.layer.0.attention.output.dense.weight | 589824 |
| encoder.layer.0.attention.output.dense.bias | 768 |
| encoder.layer.0.attention.output.LayerNorm.weight | 768 |
| encoder.layer.0.attention.output.LayerNorm.bias | 768 |
| encoder.layer.0.intermediate.dense.weight | 2359296 |
| encoder.layer.0.intermediate.dense.bias | 3072 |
| encoder.layer.0.output.dense.weight | 2359296 |
| encoder.layer.0.output.dense.bias | 768 |
| encoder.layer.0.output.LayerNorm.weight | 768 |
| encoder.layer.0.output.LayerNorm.bias | 768 |
| pooler.dense.weight | 589824 |
| pooler.dense.bias | 768 |

There are 12 such encoder layers.

So Total Trainable Params: 109482240

109482240

The model's parameter count was computed to be **109,482,240** through the code. This count closely aligns with the parameter figure of **108,432,384** as reported in the corresponding paper. The slight variance between the calculated and reported counts can potentially stem from differences in implementation details, such as parameter initialization methods and biases. The exact number of parameters from the different layers were also calculated, and have been attached in the end.

Task 3: Pretrain the selected model on the train split of 'wikitext-2-raw-v1'

The Models have been pushed to 🤗. We used the following hyperparameters:

Learning Rate = 5e-5

Epochs = 5
Batch Size: 36718
Sequence Length: 512

Models were trained using the specified hyperparameters for 5 epochs, each epoch involving batches of 36718 samples, with sequences padded or truncated to a maximum length of 512 tokens. The above hyperparameters are also inspired by the research paper.

Task 4: The Perplexity scores

| EPOCH | Mean Perplexity |
|-------|-----------------|
| 1 | 3009244.135 |
| 2 | 2351987.914 |
| 3 | 2162127.378 |
| 4 | 1900714.996 |
| 5 | 1743680.844 |

The perplexity score decreases after each epoch during training.. There are multiple reasons for this:

- Initially, as the model learns more from the training data, its ability to predict sequences improves, resulting in decreased perplexity on the test data.
- The choice of hyperparameters and the model architecture influence the convergence pattern.
- If we train on more epochs, the model starts memorizing the training data, causing overfitting. This might lead to an increase in perplexity on the test data.

Task 5: The pre-trained model has been pushed to 🤗.

Here is the link to the pushed model: [link](#)

Task 6 : The final pre-trained model has been fine-tuned on the train split of the dataset for the below tasks:

- a. Classification: SST-2
- b. Question-Answering: SQuAD

Task 7: For each of the following metrics corresponding to their particular task, the scores have been calculated on the test splits of the dataset:

- a. Classification: Accuracy, Precision, Recall, F1
Scores:

| METRIC | SCORE |
|-----------|--------|
| Recall | 1.0 |
| Precision | 0.551 |
| Accuracy | 0.551 |
| F1 | 0.7105 |

- A recall of 1.0 means that the model is capturing all the positive instances in the dataset.
- A precision of 0.551 indicates that about 55.1% of the predicted positive instances are true positives or positive sentiment.
- The accuracy of 0.551 means that the model correctly predicts the class of the instances about 55.1% of the time. In this case, it seems that the model is not performing much better than random guessing.
- An F1 score of 0.7105 is relatively good, suggesting a reasonable balance between precision and recall. However, the low precision indicates that there are still false positives that need to be addressed.

b. Question-Answering: squad_v2, F1, METEOR, BLEU, ROUGE, exact-match Scores:

| METRIC | SCORE |
|-------------------|--------|
| exact | 7.479 |
| f1 | 12.104 |
| total | 11873 |
| HasAns_exact | 14.98 |
| HasAns_f1 | 24.245 |
| HasAns_total | 5928 |
| NoAns_exact | 0 |
| NoAns_f1 | 0 |
| NoAns_total | 5945 |
| best_exact | 50.088 |
| best_exact_thresh | 0 |
| best_f1 | 50.092 |
| best_f1_thresh | 0 |

| | |
|--------------------|--------------------------|
| meteor | 0.00098 |
| bleu | 0 |
| precisions | [1.5e-05, 0.0, 0.0, 0.0] |
| brevity_penalty | 1 |
| length_ratio | 5.60 |
| translation_length | 66491 |
| reference_length | 11873 |
| rouge1 | 2.11E-05 |
| rouge2 | 0 |
| rougeL | 2.11E-05 |
| rougeLsum | 2.11E-05 |

In assessing the ROPES (Reading for Open-Ended Problem Solving) dataset, F1 score is the preferred metric, measuring a system's adeptness at applying passage knowledge to novel situations. For the evaluation of generative Question-Answering (QA) models, METEOR stands out as the optimal metric, providing a thorough evaluation of the model's performance, especially in generating responses that emulate human-like language and expression.

Task 8: Number of parameters in the model after fine-tuning are

After Fine-Tuning For Classification:

| Modules | Parameters |
|---|------------|
| <code>bert.embeddings.word_embeddings.weight</code> | 23440896 |
| <code>bert.embeddings.position_embeddings.weight</code> | 393216 |
| <code>bert.embeddings.token_type_embeddings.weight</code> | 1536 |
| <code>bert.embeddings.LayerNorm.weight</code> | 768 |
| <code>bert.embeddings.LayerNorm.bias</code> | 768 |
| <code>bert.encoder.layer.0.attention.self.query.weight</code> | 589824 |
| <code>bert.encoder.layer.0.attention.self.query.bias</code> | 768 |
| <code>bert.encoder.layer.0.attention.self.key.weight</code> | 589824 |
| <code>bert.encoder.layer.0.attention.self.key.bias</code> | 768 |
| <code>bert.encoder.layer.0.attention.self.value.weight</code> | 589824 |
| <code>bert.encoder.layer.0.attention.self.value.bias</code> | 768 |
| <code>bert.encoder.layer.0.attention.output.dense.weight</code> | 589824 |
| <code>bert.encoder.layer.0.attention.output.dense.bias</code> | 768 |
| <code>bert.encoder.layer.0.attention.output.LayerNorm.weight</code> | 768 |
| <code>bert.encoder.layer.0.attention.output.LayerNorm.bias</code> | 768 |
| <code>bert.encoder.layer.0.intermediate.dense.weight</code> | 2359296 |
| <code>bert.encoder.layer.0.intermediate.dense.bias</code> | 3072 |

| | |
|--|---------|
| bert.encoder.layer.0.output.dense.weight | 2359296 |
| bert.encoder.layer.0.output.dense.bias | 768 |
| bert.encoder.layer.0.output.LayerNorm.weight | 768 |
| bert.encoder.layer.0.output.LayerNorm.bias | 768 |
| bert.pooler.dense.weight | 589824 |
| bert.pooler.dense.bias | 768 |
| classifier.weight | 1536 |
| classifier.bias | 2 |

There were 12 such encoder layers.

Total Trainable Params: 109483778

109483778

Num of Parameters Increased by 1538 after fine-tuning for classification.

Modifications for Classification Fine-Tuning:

Two Additional Layers have been added:

classifier.weight: Matrix size of 1,536.

classifier.bias: Vector size of 2.

During fine-tuning for a classification task, an extra classifier layer (classifier.weight and classifier.bias) is added to the existing architecture of the pre-trained model. The increase in parameter count is directly attributable to the introduction of the new classifier layer, comprising classifier.weight and classifier.bias, which collectively contribute 1,538 parameters (1,536 weights + 2 biases) to the model's architecture.

After Fine-Tuning For Question Answering:

| Modules | Parameters |
|--|------------|
| bert.embeddings.word_embeddings.weight | 23440896 |
| bert.embeddings.position_embeddings.weight | 393216 |
| bert.embeddings.token_type_embeddings.weight | 1536 |
| bert.embeddings.LayerNorm.weight | 768 |
| bert.embeddings.LayerNorm.bias | 768 |
| bert.encoder.layer.0.attention.self.query.weight | 589824 |
| bert.encoder.layer.0.attention.self.query.bias | 768 |
| bert.encoder.layer.0.attention.self.key.weight | 589824 |
| bert.encoder.layer.0.attention.self.key.bias | 768 |
| bert.encoder.layer.0.attention.self.value.weight | 589824 |
| bert.encoder.layer.0.attention.self.value.bias | 768 |
| bert.encoder.layer.0.attention.output.dense.weight | 589824 |
| bert.encoder.layer.0.attention.output.dense.bias | 768 |
| bert.encoder.layer.0.attention.output.LayerNorm.weight | 768 |
| bert.encoder.layer.0.attention.output.LayerNorm.bias | 768 |

| | | | | |
|---------------|--|--|---------|--|
| | bert.encoder.layer.0.intermediate.dense.weight | | 2359296 | |
| | bert.encoder.layer.0.intermediate.dense.bias | | 3072 | |
| | bert.encoder.layer.0.output.dense.weight | | 2359296 | |
| | bert.encoder.layer.0.output.dense.bias | | 768 | |
| | bert.encoder.layer.0.output.LayerNorm.weight | | 768 | |
| | bert.encoder.layer.0.output.LayerNorm.bias | | 768 | |
| | qa_outputs.weight | | 1536 | |
| | qa_outputs.bias | | 2 | |
| +-----+-----+ | | | | |

There were 12 such encoder layers.
Total Trainable Params: 108893186
108893186

Changes in Layer Parameters after QA Fine-tuning

The Original Layers have been changed to new layers:

Original Model Layers:

pooler.dense.weight: Matrix size of 589,824 (e.g., for a BERT-based model).
pooler.dense.bias: Vector size of 768.

Modifications for Question Answer Fine-Tuning:

Adjusted Layers:

qa_outputs.weight: Matrix size of 1,536.
qa_outputs.bias: Vector size of 2.

The fine-tuning involves replacing the pre-trained model's original final layers (pooler.dense.weight and pooler.dense.bias) with fine-tuned layers (qa_outputs.weight and qa_outputs.bias) optimized for QA tasks. The decrease in parameter count is primarily due to the reduction in the dimensions of the final layers from the original size (589,824 weights and 768 biases) to the tailored QA-specific layers (1,536 weights and 2 biases).

Task 9: The fine-tuned model has been pushed to 😊 : [Fine-Tuned Classification](#), [Fined-Tune-QA](#)

Comments and Discussion:

Performance

The model's performance is poor and the scores from the performance metrics is not up to the mark as:

- The model's Exact Match (EM) score of 7.48% indicates a relatively low accuracy in providing exact answers to questions in the dataset.
- The F1 score of 12.10% suggests a modest balance between precision and recall, but there is room for improvement.

- The model performs better on questions with answers (HasAns) with an EM score of 14.98% and an F1 score of 24.24%. This indicates some ability to answer questions when information is available.
- The model's performance on questions with no answers (NoAns) is notably lower, with an EM and F1 score of 0.00%. This suggests challenges in identifying unanswerable questions.
- The best exact match and F1 scores are relatively higher at 50.09%, indicating that the model can achieve better performance under certain conditions or thresholds.
- The METEOR score of 0.00098 is quite low, suggesting that the model may struggle with generating responses that closely align with human-like language and expression.
- The BLEU score of 0.00 indicates minimal overlap between the model's generated text and reference text.
- The ROUGE scores are also very low, suggesting limited similarity between the model's outputs and reference text.

Work Distribution

| Name | Contribution |
|---------|--|
| Aditi | Calculated number of parameters of selected model, Helped in Pre-training |
| Anavart | Fine-Tuned for Classification |
| Ayush | Fine-Tuned for Question & Answering, Calculated Metric Scores for Question-Answering, Performance Inference, Documentation |
| Mann | Pre-Training the BERT Model using MLM and NSP, Pushed model to HuggingFace, Calculated perplexity, Documentation |
| Pinki | Pre-Training the Bert Model using MLM and NSP |
| Ronak | Fine-Tuned for Question & Answering, Calculated metrics for the Classification task |
| Sakshi | Fine-Tuned for Question & Answering, Documentation |
| Vedang | Pre Training the BERT Model, Fine-tuning for classification, Error analysis and debugging, Calculated perplexity |